

AWS-Based Fraud Detection Demo

Objective

Build a lightweight end-to-end skeleton of a fraud detection ML system using AWS services to highlight design bottlenecks that MLOps solves.

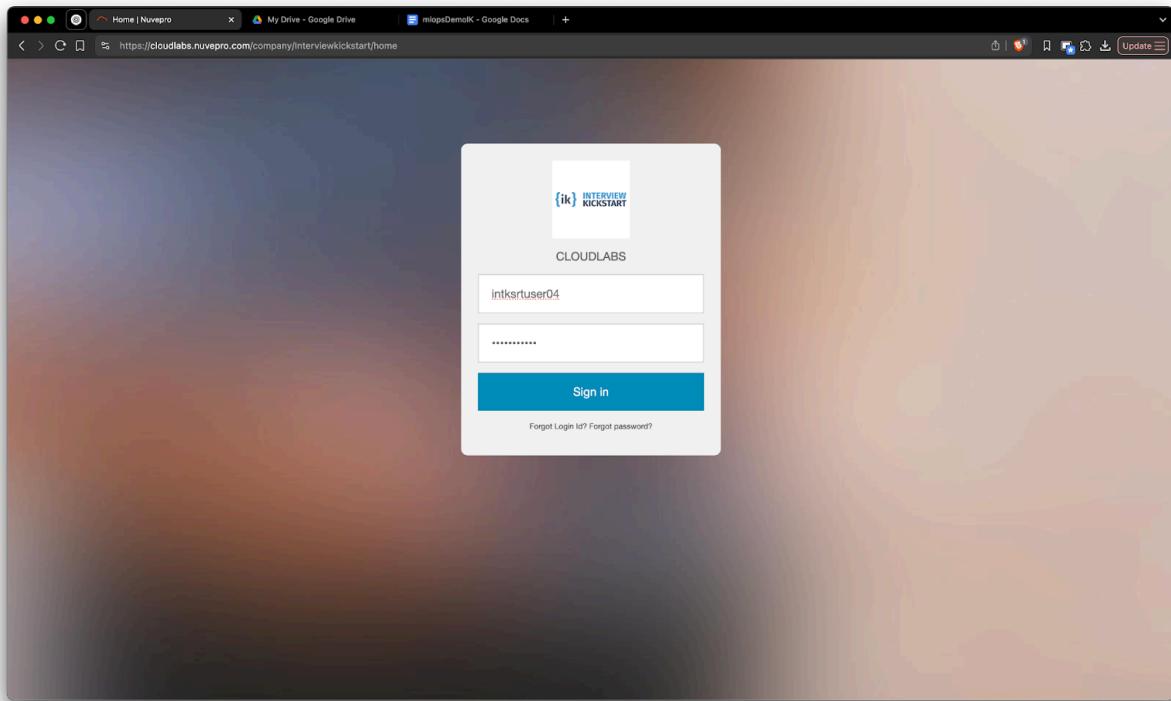
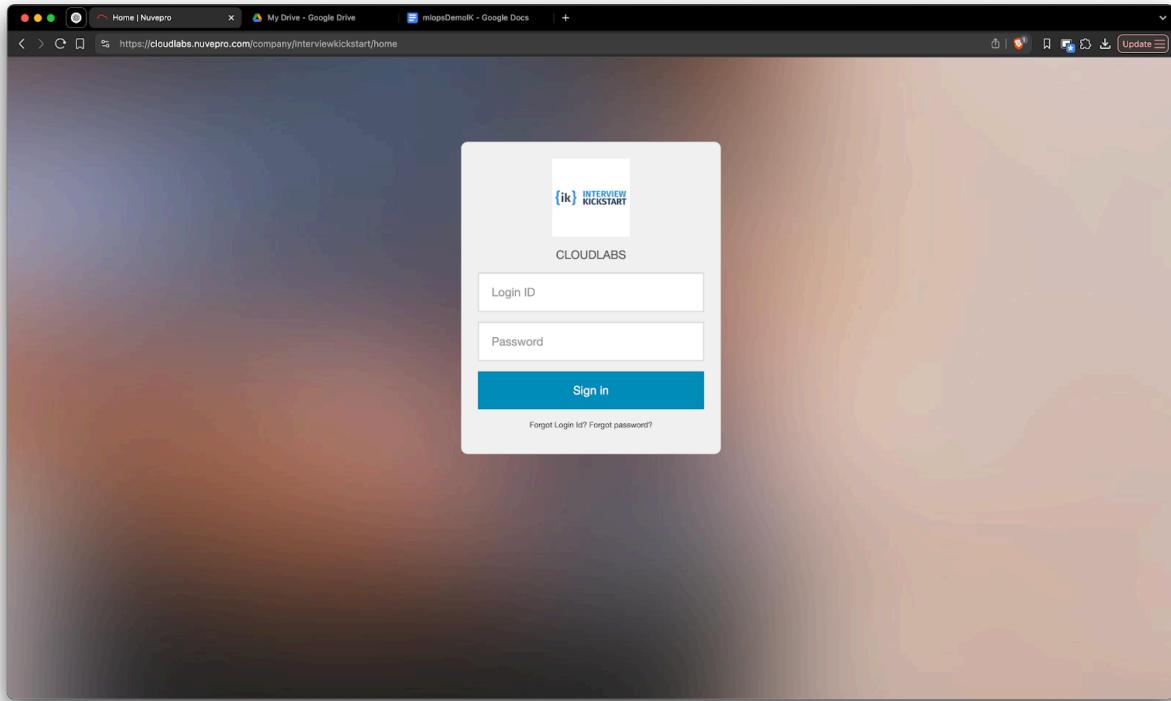
Prerequisites

- AWS account with billing enabled
 - IAM user with admin or required permissions
 - AWS CLI installed and configured
 - Python 3.8+ environment (EC2 or local)
 - Kaggle account and API token (kaggle.json)
-

Step 0: Access AWS Console via Nuvepro

Goal: Begin from the Nuvepro-provided cloud lab environment

1. Log in to your Nuvepro account.
2. Navigate to the lab dashboard.
3. Click on the assigned AWS environment.
4. Use provided credentials or direct link to access the AWS Console.



Dashboard | Nuvepro My Drive - Google Drive mlopsDemoK - Google Docs

https://cloudlabs.nuvepro.com/dashboard

Home ▾

My Labs (1)

Welcome Intksrt User04 !

AWS Lab

#1904977

Amazon Web Services(AWS) account and access to all resources available in AWS

Tools : None
Allocated : Unlimited
Created On : 29-Sep-2023 03:13:52 AM
Expires On : N/A

[View Lab](#) [Stop-Complete](#)

<https://cloudlabs.nuvepro.com/subscriptions/launch?id=1904977>

Subscription Details | Nuvepro My Drive - Google Drive mlopsDemoK - Google Docs

https://cloudlabs.nuvepro.com/subscriptions/launch?id=1904977

Home ▾

Lab Control Panel

AWS Lab

[Start](#)

Latest Status

[Get Usage](#) [Trends...](#)

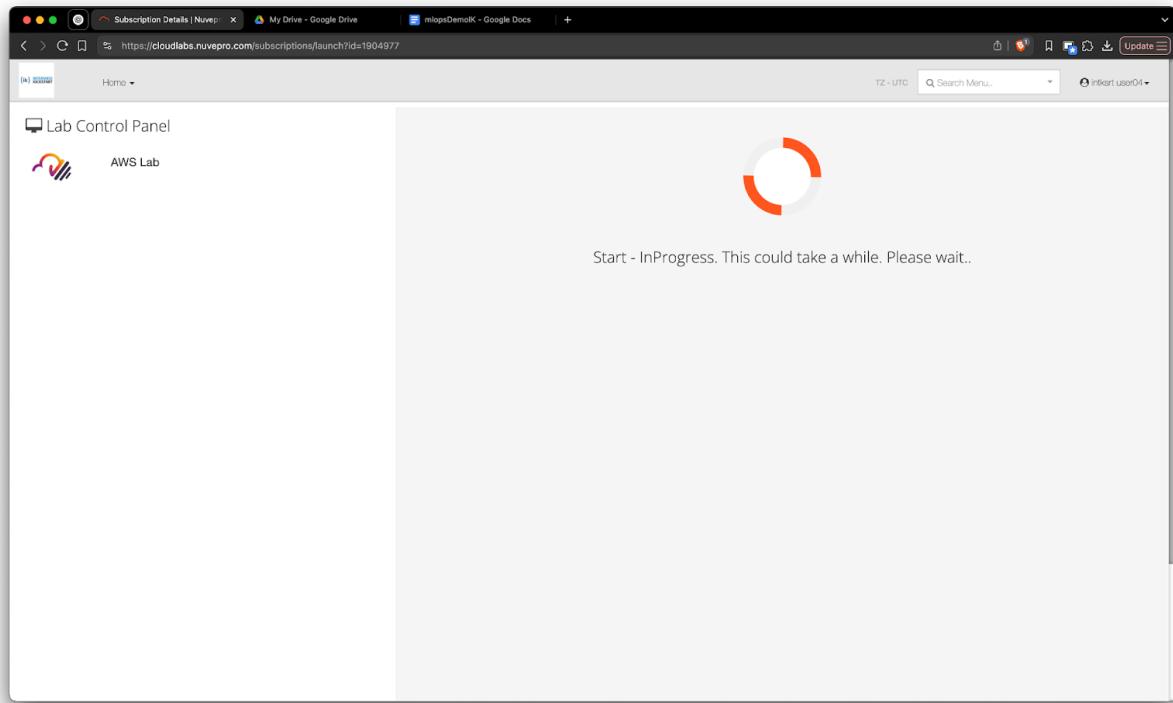
Note : There is a delay of 24 to 48 hours for the usage cost to be generated by the Cloud platform. Please note that the cost shown for this lab may not be complete

Usage Trends Events Feedback

Owner - intksrt04User@nuvelabs.com

 **⚠**

Lab not started. Please retry.



The screenshot shows the 'Lab Control Panel' for an 'AWS Lab'. The interface includes a sidebar with 'Access Details' and a main panel with tabs for 'Usage Trends', 'Events', and 'Feedback'. A red 'Stop' button is visible. The 'More Details' section is expanded, showing links for 'Policies', 'Instructions', 'Other Details', and 'Quota'. A 'Jump to Console' button is present.

Access Details

loginId	inkartuser04
loginPassword	*****
isReadyToUse	true
userName	inkartuser04
registeredMailId	kawalsi457@nuvelabs.com
userId	467017
parentId	o-p9t2rnkmdo
accessId	AKAAXRMNPSV74E63DDVT
userList	[]
loginUser	inkartuser04
accessKey	gloGB1S0WwJW1BgvGSEvGe2j81EXUpOH r+3W
tenantId	228
id	160003
state	ACTIVE
arn	518377149823
password	*****

Usage Cost

The screenshot shows the same 'Lab Control Panel' interface. A green 'Action Completed' banner is at the top right. A modal dialog titled 'Connection Details' contains an 'OPEN CONSOLE' button and a warning message: 'Warning: Click on the above button to access the AWS lab. All the resources created in the AWS lab will be deleted after the duration as specified in the user guide. please plan your labs accordingly.' The 'More Details' section is partially visible behind the modal.

Action Completed

Connection Details

OPEN CONSOLE

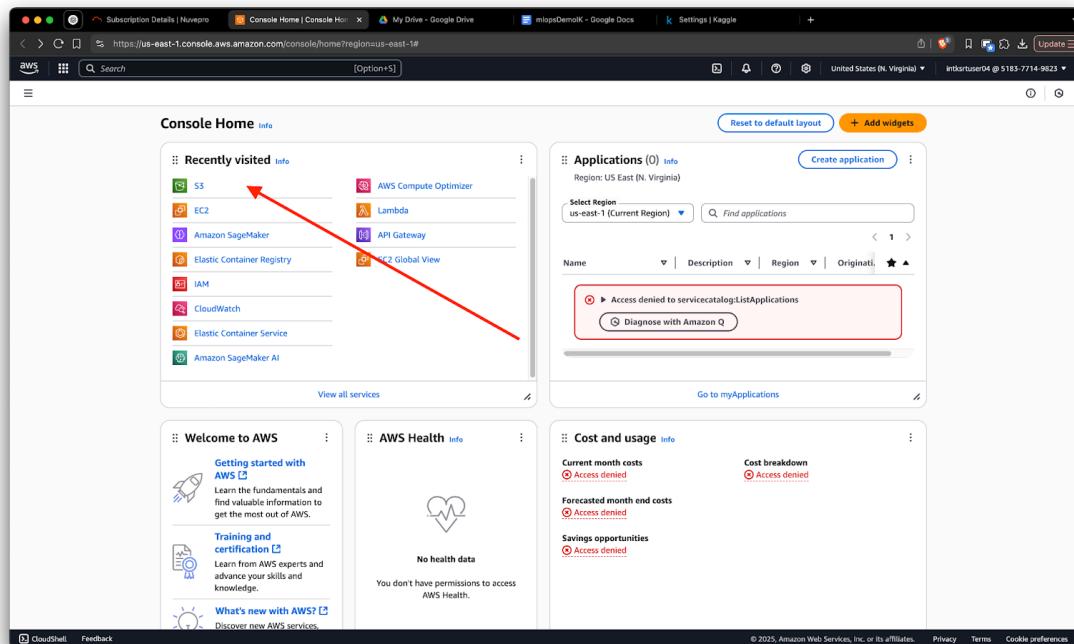
Warning: Click on the above button to access the AWS lab. All the resources created in the AWS lab will be deleted after the duration as specified in the user guide. please plan your labs accordingly.

Nuvepro offers sandboxed environments for training and demos. Ensure your session remains active.

Step 1: Setup S3 Bucket

Goal: Store raw data and simulate ingestion

1. Go to AWS Console > S3
2. Create bucket: fraud-demo-data-yourname
3. Enable **versioning**
4. **Block all public access**



Amazon S3

General purpose buckets

Storage lens provides visibility into storage usage and activity trends. Metrics don't include directory buckets. [Learn more](#)

[View Storage Lens dashboard](#)

General purpose buckets (10) All AWS Regions

Buckets are containers for data stored in S3.

[Find buckets by name](#)

Name	AWS Region	IAM Access Analyzer	Creation date
ashok-ovnaelvianlaud	US East (N. Virginia) us-east-1	View analyzer for us-east-1	December 13, 2024, 08:16:55 (UTC+05:30)
ashok-s3	US East (N. Virginia) us-east-1	View analyzer for us-east-1	June 22, 2024, 00:10:52 (UTC+05:30)
ashok-ik-test	US East (N. Virginia) us-east-1	View analyzer for us-east-1	June 25, 2024, 07:41:54 (UTC+05:30)
churn-classifier-data-aj	US East (N. Virginia) us-east-1	View analyzer for us-east-1	November 1, 2024, 07:14:55 (UTC+05:30)
dm-assign-ik	US East (N. Virginia) us-east-1	View analyzer for us-east-1	September 21, 2024, 04:36:40 (UTC+05:30)
dm-demo-ik	US East (N. Virginia) us-east-1	View analyzer for us-east-1	September 15, 2024, 19:33:07 (UTC+05:30)
dm-demo-ik-copy	US East (N. Virginia) us-east-1	View analyzer for us-east-1	September 15, 2024, 19:33:25 (UTC+05:30)
mlopsashok	US East (N. Virginia) us-east-1	View analyzer for us-east-1	December 9, 2024, 02:18:54 (UTC+05:30)
nm-lambda-demo-2024-05-12	US East (N. Virginia) us-east-1	View analyzer for us-east-1	May 12, 2024, 23:20:49 (UTC+05:30)
sentiment-classifier-data-aj	US East (N. Virginia) us-east-1	View analyzer for us-east-1	October 27, 2024, 04:50:55 (UTC+05:30)

[Create bucket](#)

Create bucket Info

Buckets are containers for data stored in S3.

General configuration

AWS Region: US East (N. Virginia) us-east-1

Bucket type: [Info](#) General purpose Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

Directory Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket name: [Info](#) **fraud-demo-data**

Bucket names must be 3 to 63 characters and unique within the global namespace. Bucket names must also begin and end with a letter or number. Valid characters are a-z, 0-9, periods (.), and hyphens (-). [Learn More](#)

Copy settings from existing bucket - optional: Only the bucket settings in the following configuration are copied. [Choose bucket](#)

Format: s3://[bucket]/prefix

Object Ownership Info

Control ownership of objects written to this bucket from other AWS accounts and the use of access control lists (ACLs). Object ownership determines who can specify access to objects.

ACLs disabled (recommended) All objects in this bucket are owned by this account. Access to this bucket and its objects is specified using only policies.

ACLs enabled Objects in this bucket can be owned by other AWS accounts. Access to this bucket and its objects can be specified using ACLs.

Object Ownership: Bucket owner enforced

Block Public Access settings for this bucket

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, access point policies, or alt. In order to ensure that public access to this bucket and its objects is blocked, turn on Block all public access. These settings apply only to this bucket and its access points. AWS recommends that you turn on Block all public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require some level of public access to this bucket or objects within, you can customize the individual settings below to suit your specific storage use cases. [Learn more](#)

This bucket is your data lake. Versioning ensures traceability and reproducibility in ML pipelines.

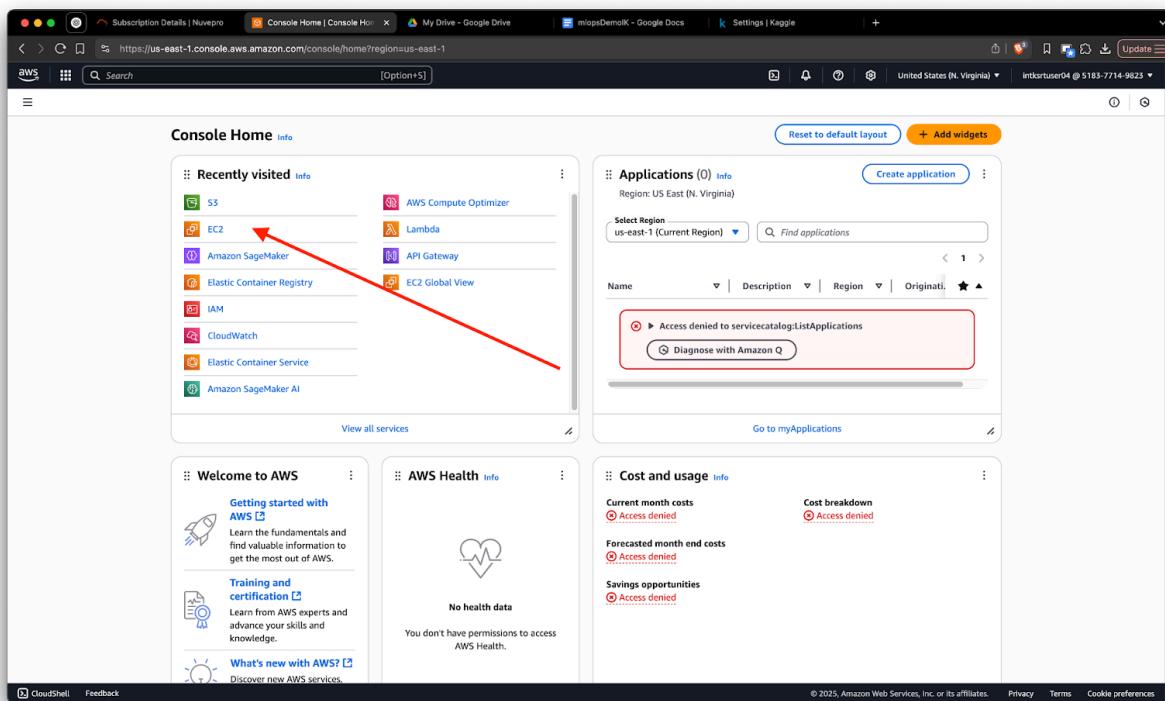
Step 2: Launch EC2 Instance (Compute Node)

Goal: Act as the data ingestion + training server

1. Launch an EC2 instance with **Ubuntu 22.04**
2. Attach IAM Role with these permissions:
 - AmazonS3FullAccess
 - AmazonSageMakerFullAccess
 - CloudWatchFullAccess
 - AmazonSSMManagedInstanceCore
3. Add tag: Purpose=fraud-demo

Access options: - **SSH:** Create/select key pair, allow port 22 - **Session**

Manager: Ensure IAM role has SSM access and instance is in public subnet



Screenshot of the AWS EC2 Dashboard (https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#Home):

- Resources:** Shows 0 Instances (running), 0 Auto Scaling Groups, 0 Capacity Reservations, 0 Dedicated Hosts, 0 Elastic IPs, 0 Instances, 8 Key pairs, 5 Load balancers, 0 Placement groups, 25 Security groups, 0 Snapshots, 0 Volumes.
- Launch instance:** A red arrow points to the "Launch instance" button.
- Service health:** An error occurred retrieving service health information.
- Zones:** Lists availability zones: us-east-1a (use1-az2), us-east-1b (use1-az4), us-east-1c (use1-az6), us-east-1d (use1-az1), us-east-1e (use1-az3), us-east-1f (use1-az5).
- Account attributes:** Default VPC (vpc-0b38dd1476f9963d0), Settings (Data protection and security, Allowed AMIs, Zones, EC2 Serial Console, Default credit specification, EC2 console preferences).
- Explore AWS:** Amazon GuardDuty Malware Protection, Get Up to 40% Better Price Performance, Save up to 90% on EC2 with Spot Instances.
- Additional information:** Get started walkthroughs, Getting started guide, Documentation.

Screenshot of the "Launch an instance" wizard (https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#LaunchInstances):

- Name and tags:** Name field contains "mlops-mydemo".
- Application and OS Images (Amazon Machine Image):** An AMI is selected: "Ubuntu Server 22.04 LTS (HVM), SSD Volume Type" (ami-0f9de6e2d2f067fc). A red arrow points to the "Launch instance" button.
- Summary:** Number of instances: 1. Includes sections for Software Image (AMI), Virtual server type (instance type) (t2.micro), Firewall (security group) (New security group), Storage (volumes) (1 volume(s) - 8 GiB), and a note about the Free tier.

Screenshot of the AWS EC2 Launch Instance wizard, Step 3: Configure Instance Details.

Instance type: t2.micro

Key pair (login): ik-pj-peM

Network settings:

- Subnet: No preference (Default subnet in any availability zone)
- Auto-assign public IP: Enabled
- Firewall (security groups): Create security group

Create key pair:

Key pair name: Enter key pair name

Key pair type: RSA (selected)

Private key file format: pem (selected)

Warning: When prompted, store the private key in a secure and accessible location on your computer. You will need it later to connect to your instance. Learn more

Software Image (AMI): Canonical, Ubuntu, 22.04, amd64

Virtual server type (instance type): t2.micro

Storage (volumes): 1 volume(s) - 8 GiB

Summary: Number of instances: 1

Launch instance button

Screenshot of the AWS EC2 Launch Instance wizard, Step 4: Review and Launch.

Success: Successfully initiated launch of instance (i-0606147e4d12ds1e)

Next Steps:

- Launch log
- What would you like to do next with this instance, for example "create alarm" or "create backup?"

Next Steps (Card View):

- Create billing and free tier usage alerts**: Create billing alerts
- Connect to your instance**: Connect to instance, Learn more
- Connect an RDS database**: Connect an RDS database, Create a new RDS database, Learn more
- Create EBS snapshot policy**: Create EBS snapshot policy
- Manage detailed monitoring**: Manage detailed monitoring
- Create Load Balancer**: Create Load Balancer
- Create AWS budget**: Create AWS budget
- Manage CloudWatch alarms**: Manage CloudWatch alarms

View all instances button

Screenshot of the AWS EC2 Instances page showing a single instance named "mlops-mydemo".

Instances (1) Info

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 IP
mlops-mydemo	i-0606147e4d12d5a1e	Running	t2.micro	Initializing		us-east-1a	ec2-3-85-62-155.compute...	3.85.62.155

Select an instance

Screenshot of the AWS EC2 Instance details page for the instance "i-0606147e4d12d5a1e".

Instance summary for i-0606147e4d12d5a1e (mlops-mydemo)

Instance ID	i-0606147e4d12d5a1e	Public IPv4 address	3.85.62.155 open address
IPv6 address	-	Instance state	Running
Hostname type	IP name: ip-172-31-82-254.ec2.internal	Private IP DNS name (IPv4 only)	ip-172-31-82-254.ec2.internal
Answer private resource DNS name	IPv4 (A)	Instance type	t2.micro
Auto-assigned IP address	3.85.62.155 [Public IP]	VPC ID	vpc-0b38dd1476f9963d0
IAM Role	-	Subnet ID	subnet-0c253255d912892f4
IMDSv2	Required	Instance ARN	arn:aws:ec2:us-east-1:518377149823:instance/i-0606147e4d12d5a1e
Operator	-		

Details **Status and alarms** **Monitoring** **Security** **Networking** **Storage** **Tags**

Instance details

AMI ID	i-0f9de6e2d2f067fc4	Monitoring	disabled
AMI name	ubuntu/images/hvm-ssd/ubuntu-jammy-22.04-amd64-server-20250305	Allowed image	-

Platform details

Linux/UNIX

Termination protection

Disabled

The screenshot shows the AWS EC2 Instances details page for an instance named 'i-0606147e4d12d5a1e'. The left sidebar navigation includes 'EC2' (selected), 'Dashboard', 'EC2 Global View', 'Events', 'Instances' (selected), 'Instance Types', 'Launch Templates', 'Spot Requests', 'Savings Plans', 'Reserved Instances', 'Dedicated Hosts', 'Capacity Reservations', 'Images' (AMIs, AMI Catalog), 'Elastic Block Store' (Volumes, Snapshots, Lifecycle Manager), 'Network & Security' (Security Groups, Elastic IPs, Placement Groups, Key Pairs, Network Interfaces), and 'Load Balancing' (Load Balancers, Target Groups). The main content area displays the 'Instance summary' for the selected instance. The 'Details' tab is active. Key details include:

- Instance ID:** i-0606147e4d12d5a1e
- IP address:** -
- Hostname type:** IP name: ip-172-31-82-254.ec2.internal
- Answer private resource DNS name:** IPv4 (A)
- Auto-assigned IP address:** 3.85.62.155 [Public IP]
- IAM Role:** -
- IMDSv2:** Required
- Operator:** -
- Public IPv4 address:** 3.85.62.155 | open address
- Instance state:** Running
- Private IP DNS name (IPv4 only):** ip-172-31-82-254.ec2.internal
- Instance type:** t2.micro
- VPC ID:** vpc-0b321476f9963d0
- Subnet ID:** subnet-0c253255d912892f4
- Instance ARN:** arnaws:ec2:us-east-1:518377149823:instance/i-0606147e4d12d5a1e

On the right side, there are sections for 'Private IPv4 addresses' (172.31.82.254), 'Public IPv4 DNS' (ec2-3-85-62-155.compute-1.amazonaws.com), 'Elastic IP addresses' (none), 'AWS Compute Optimizer finding' (none), 'Auto Scaling Group name' (none), and 'Managed' (false). The top right of the main content area has a 'Connect' button, which is highlighted by a red arrow.

The screenshot shows the AWS EC2 Connect interface. At the top, there are tabs for 'EC2 Instance Connect', 'Session Manager', 'SSH client' (which is selected), and 'EC2 serial console'. Below the tabs, the 'Instance ID' is listed as 'i-0606147e4d12d5a1e (mllops-mydemo)'. A numbered list of steps for connecting via SSH is provided:

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is `ikdemopj.pem`.
3. Run this command, if necessary, to ensure your key is not publicly viewable.
4. Connect to your instance using its Public DNS:

Below the steps, an 'Example:' section shows the command: `ssh -i "ikdemopj.pem" ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com`. A note at the bottom states: 'Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.'

The screenshot shows the AWS EC2 Connect interface with the 'SSH client' tab selected. The 'Instance ID' is 'i-0606147e4d12d5a1e (mllops-mydemo)'. The same numbered connection steps are present. A message 'Command copied' is displayed next to the fourth step's button. The terminal window on the right shows the copied command: `ssh -i "ikdemopj.pem" ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com`. The terminal also displays system information and a note about security maintenance.

```
pranjul@mbpArxiv temporary % chmod 400 "ikdemopj.pem"
pranjul@mbpArxiv temporary % ssh -i "ikdemopj.pem" ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com
The authenticity of host 'ec2-3-85-62-155.compute-1.amazonaws.com (3.85.62.155)' can't be established.
ED25519 key fingerprint is SHA256:Utd0i1HMKxtk33Lir>xRf0AdNQqrHLmDsJmzYDK8Hi0.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[Fingerprint])? yes
Warning: Permanently added 'ec2-3-85-62-155.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
Welcome to Ubuntu 22.04.5 LTS (GNU/Linux 6.8.0-1824-aws x86_64)

System information as of Tue May 13 18:57:30 UTC 2025
System load: 0.16      Processes:          107
Usage of /: 21.8% of 7.57GB   Users logged in:        0
Memory usage: 22%           IPv4 address for eth0: 172.31.82.254
Swap usage: 0%
```

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See <https://ubuntu.com/esm> or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

```
ubuntu@ip-172-31-82-254:~$
```

The screenshot shows the AWS IAM Dashboard. On the left, there's a sidebar with 'Access management' expanded, showing 'User groups', 'Users', 'Roles' (which has a red arrow pointing to it), 'Policies', 'Identity providers', 'Account settings', and 'Root access management'. Below that are sections for 'Access reports', 'Access Analyzer', 'Credential report', 'Organization activity', 'Service control policies', and 'Resource control policies'. At the bottom of the sidebar are links for 'IAM Identity Center' and 'AWS Organizations'. The main area is the 'IAM Dashboard' with sections for 'IAM resources' (User groups: 6, Users: 1, Roles: 13, Policies: 0, Identity providers: 0) and 'AWS Account' (Account ID: 518577149823, Account Alias: Create, Sign-in URL: https://518577149823.signin.aws.amazon.com/console). There's also a 'What's new' section with a list of recent updates and a 'Tools' section with a 'Policy simulator'. At the bottom of the dashboard are links for 'CloudShell', 'Feedback', and copyright information.

Screenshot of the AWS IAM 'Create role' wizard Step 1: Select trusted entity.

The 'Trusted entity type' section shows the following options:

- AWS service**: Allow AWS services like EC2, Lambda, or others to perform actions in this account.
- AWS account**: Allow entities in other AWS accounts belonging to you or a 3rd party to perform actions in this account.
- Web Identity**: Allows users federated by the specified external web identity provider to assume this role to perform actions in this account.
- SAML 2.0 federation**: Allow users federated with SAML 2.0 from a corporate directory to perform actions in this account.
- Custom trust policy**: Create a custom trust policy to enable others to perform actions in this account.

The 'Use case' section shows the following options under 'Service or use case':

- EC2**: Allows EC2 Instances to call AWS services on your behalf.
- EC2 Role for AWS Systems Manager**: Allows EC2 Instances to call AWS services like CloudWatch and Systems Manager on your behalf.
- EC2 Spot Fleet Role**: Allows EC2 Spot Fleet to request and terminate Spot Instances on your behalf.
- EC2 - Spot Fleet Auto Scaling**: Allows Auto Scaling to access and update EC2 spot fleets on your behalf.
- EC2 - Spot Fleet Logging**: Allows EC2 Spot Instances to attach tags to the launched instances on your behalf.
- EC2 - Spot Instances**: Allows EC2 Spot Instances to launch and manage spot instances on your behalf.
- EC2 - Spot Fleet**: Allows EC2 Spot Fleet to launch and manage spot fleet instances on your behalf.
- EC2 - Scheduled Instances**: Allows EC2 Scheduled Instances to manage instances on your behalf.

Buttons at the bottom: **Cancel**, **Next**.

Screenshot of the AWS IAM 'Create role' wizard Step 2: Add permissions.

The 'Permissions policies' section shows the following policies:

- CloudWatchFullAccess**: AWS managed
- CloudWatchFullAccessV2**: AWS managed

The 'Permissions' sidebar provides information about policies:

You manage access in AWS by creating policies and attaching them to IAM identities (users, groups of users, or roles) or AWS resources. A policy is an object in AWS that, when associated with an identity or resource, defines their permissions.

Was this content helpful? **Yes** **No**

Learn more **Creating IAM policies**

Buttons at the bottom: **Cancel**, **Previous**, **Next**.

The screenshot shows the AWS IAM 'Create role' wizard at Step 2: 'Add permissions'. The left sidebar shows steps: Step 1 (Select trusted entity), Step 2 (Add permissions, which is selected), and Step 3 (Name, review, and create). The main area is titled 'Add permissions' and shows a search bar for 'Permissions policies (2/1046)'. A table lists two policies: 'AmazonSageMakerFullAccess' and 'AmazonSageMakerFullAccess'. The right sidebar contains a 'Permissions' section with a detailed description of what policies are for, and a 'Was this content helpful?' poll.

This screenshot is identical to the one above, showing the 'Add permissions' step of the 'Create role' wizard. It displays the same search results for AWS managed policies and the same right-hand sidebar with the 'Permissions' information and helpfulness poll.

Subscription Details | Nuvepro Launch an instance | EC2 | us-east-1 Create role | IAM | Global My Drive - Google Drive mlopsDem0K - Google Docs Settings | Kaggle

https://us-east-1.console.aws.amazon.com/iam/home?region=us-east-1#/roles/create?trustedEntityType=AWS_SERVICE&selectedService=EC2&selectedUseCase=EC2&policies=arn%3aws%3...

IAM Roles Create role [Options]

Step 1 Select trusted entity Step 2 Add permissions Step 3 Name, review, and create

Name, review, and create

Role details

Role name Enter a meaningful name to identify this role.
IdentityIops Maximum 64 characters. Use alphanumeric and +/-_ characters.

Description Add a brief explanation for this role.
Allows EC2 instances to call AWS services on your behalf. Maximum 1000 characters. Use letters (A-Z and a-z), numbers (0-9), tabs, new lines, or any of the following characters: _-_.@!#\$%^&`~"-

Step 1: Select trusted entities

Trust policy

```
1: {  
2: "Version": "2012-10-17",  
3: "Statement": [  
4: {  
5: "Effect": "Allow",  
6: "Action": "sts:AssumeRole",  
7: "Principal": {  
8: "Service": [  
9: "ec2.amazonaws.com"  
10: ]  
11: }  
12: }  
13: ]  
14: }  
15: ]  
16: }
```

Step 2: Add permissions

Permissions policy summary

Policy name	Type	Attached as
AmazonSNSFullAccess	AWS managed	Permissions policy
AmazonSageMakerFullAccess	AWS managed	Permissions policy
CloudWatchFullAccess	AWS managed	Permissions policy

Step 3: Add tags

Add tags - optional Tags are key-value pairs that you can add to AWS resources to help identify, organize, or search for resources. No tags associated with the resource.

Add new tag You can add up to 50 more tags.

Cancel Previous Create role Next Cancel Previous Create role Next

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookies preferences

Subscription Details | Nuvepro Launch an instance | EC2 | us-east-1 Create role | IAM | Global My Drive - Google Drive mlopsDem0K - Google Docs Settings | Kaggle

https://us-east-1.console.aws.amazon.com/iam/home?region=us-east-1#/roles/create?trustedEntityType=AWS_SERVICE&selectedService=EC2&selectedUseCase=EC2&policies=arn%3aws%3...

IAM Roles Create role [Options]

Step 1 Select trusted entity Step 2 Add permissions Step 3 Name, review, and create

Name, review, and create

Role details

Role name Enter a meaningful name to identify this role.
IdentityIops Maximum 64 characters. Use alphanumeric and +/-_ characters.

Description Add a brief explanation for this role.
Allows EC2 instances to call AWS services on your behalf. Maximum 1000 characters. Use letters (A-Z and a-z), numbers (0-9), tabs, new lines, or any of the following characters: _-_.@!#\$%^&`~"-

Step 1: Select trusted entities

Trust policy

```
1: {  
2: "Version": "2012-10-17",  
3: "Statement": [  
4: {  
5: "Effect": "Allow",  
6: "Action": "sts:AssumeRole",  
7: "Principal": {  
8: "Service": [  
9: "ec2.amazonaws.com"  
10: ]  
11: }  
12: }  
13: ]  
14: }  
15: ]  
16: }
```

Step 2: Add permissions

Permissions policy summary

Policy name	Type	Attached as
AmazonSNSFullAccess	AWS managed	Permissions policy
AmazonSageMakerFullAccess	AWS managed	Permissions policy
CloudWatchFullAccess	AWS managed	Permissions policy

Step 3: Add tags

Add tags - optional Tags are key-value pairs that you can add to AWS resources to help identify, organize, or search for resources. No tags associated with the resource.

Add new tag You can add up to 50 more tags.

Cancel Previous Create role Next Cancel Previous Create role Next

© 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookies preferences

Screenshot of the AWS EC2 Instances page showing a single running t2.micro instance.

The page title is "Instances (1) | EC2 | AWS Management Console".

The instance details table:

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 IP	Elastic IP	IPv6 IPs
<input type="checkbox"/>	mlops-mydemo	i-0606147e4d12d5a1e	Running	t2.micro	2/2 checks pass	View alarms +	us-east-1a	ec2-3-85-62-155.com...	3.85.62.155	-	-

Left sidebar navigation:

- EC2 > Instances
- Dashboard
- EC2 Global View
- Events
- Instances
 - Instance Types
 - Launch Templates
 - Spot Requests
 - Savings Plans
 - Reserved Instances
 - Dedicated Hosts
 - Capacity Reservations
- Images
 - AMIs
 - AMI Catalog
- Elastic Block Store
 - Volumes
 - Snapshots
 - Lifecycle Manager
- Network & Security
 - Security Groups
 - Elastic IPs
 - Placement Groups
 - Key Pairs
 - Network Interfaces
- Load Balancing
 - Load Balancers
 - Target Groups
 - Trust Stores
- Auto Scaling
 - Auto Scaling Groups

Bottom footer:

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Instance summary for i-0606147e4d12d5a1e (mlops-mydemo)

Details **Status and alarms** **Monitoring** **Security** **Networking** **Storage** **Tags**

Actions

- Connect
- Instance state
- Actions ▾
- Private IPv4 addresses
- Public IPv4 DNS
- Public IPv6 DNS
- Change security groups
- Get Windows password
- Modify IAM role
- elastic IP addresses
- Connect
- Manage instance state
- Instance settings
- Networking
- Security
- Image and templates
- Monitor and troubleshoot

Modify IAM role

Attach an IAM role to your instance.

Instance ID: i-0606147e4d12d5a1e (mlops-mydemo)

IAM role: Select an IAM role to attach to your instance or create a new role if you haven't created any. The role you select replaces any roles that are currently attached to your instance.

Create new IAM role

Update IAM role

idemoplops

- No IAM Role
- Choose this option to detach an IAM role.
- idemoplops** (arn:aws:iam::518377149823:instance-profile/idemoplops)
- s3-from-ec2** (arn:aws:iam::518377149823:instance-profile/s3-from-ec2)
- s3-access-from-ec2** (arn:aws:iam::518377149823:instance-profile/s3_access_from_ec2)
- s3-access-from_ec2** (arn:aws:iam::518377149823:instance-profile/s3_access_from_ec2)
- s3-access_from_ec2_a** (arn:aws:iam::518377149823:instance-profile/s3_access_from_ec2_a)

IAM roles provide secure access to services without hardcoding credentials.

Step 3: Configure EC2 Environment, Download and Upload Dataset to S3

```
→sudo apt update && sudo apt install -y python3-pip unzip  
→pip3 install kaggle boto3 pandas watchtower  
→mkdir ~/.kaggle
```

If kaggle not found:

```
→pip3 install kaggle  
→export PATH="$PATH:~/local/bin"  
→echo 'export PATH="$PATH:~/local/bin"' >> ~/.bashrc  
→source ~/.bashrc
```

Upload kaggle.json: - **SSH**: scp -i your-key.pem kaggle.json
ubuntu@<EC2-IP>:~/kaggle

(option) Session Manager: Use “Upload File” in AWS Console

```
→chmod 600 ~/kaggle/kaggle.json
```

```
temporary — ubuntu@ip-172-31-82-254: ~ — ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com — 123x44
~/Documents/temporary — ubuntu@ip-172-31-82-254: ~ — ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com

[pranjal@mbpArxiv temporary % chmod 400 "ikdemopj.pem"
[pranjal@mbpArxiv temporary % ssh -i "ikdemopj.pem" ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com
The authenticity of host 'ec2-3-85-62-155.compute-1.amazonaws.com (3.85.62.155)' can't be established.
ED25519 key fingerprint is SHA256:Ut0i1HMHxi83Lir+xrf6AdMQqxRHLmDsJm73YDK8Hi0.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-3-85-62-155.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
Welcome to Ubuntu 22.04.5 LTS (GNU/Linux 6.8.0-1024-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/pro

System information as of Tue May 13 18:57:30 UTC 2025

System load: 0.16          Processes:      107
Usage of /: 21.8% of 7.57GB   Users logged in:    0
Memory usage: 22%           IPv4 address for eth0: 172.31.82.254
Swap usage:  0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-82-254:~$ sudo apt update && sudo apt install -y python3-pip unzip
```

```
Running kernel seems to be up-to-date.

Restarting services...
Service restarts being deferred:
/etc/needrestart/restart.d/dbus.service
systemctl restart networkd-dispatcher.service
systemctl restart polkit.service
systemctl restart unattended-upgrades.service

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
[ubuntu@ip-172-31-82-254: ~] ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com
~/Documents/temporary ~ ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com + [ ]
Running kernel seems to be up-to-date.

Restarting services...
Service restarts being deferred:
/etc/needrestart/restart.d/dbus.service
systemctl restart networkd-dispatcher.service
systemctl restart polkit.service
systemctl restart unattended-upgrades.service

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
[ubuntu@ip-172-31-82-254: ~] pip3 install kaggle boto3 pandas
Defaulting to user installation because normal site-packages is not writeable
Collecting kaggle
  Downloading kaggle-1.7.4.5-py3-none-any.whl (181 kB)
    181.2/181.2 KB 4.3 MB/s eta 0:00:00
Collecting boto3
  Downloading boto3-1.38.14-py3-none-any.whl (139 kB)
    139.9/139.9 KB 20.5 MB/s eta 0:00:00
Collecting pandas
  Downloading pandas-2.2.3-cp310-cp310-manylinux_2_17_x86_64_manylinux2014_x86_64.whl (13.1 MB)
    13.1/13.1 MB 79.4 MB/s eta 0:00:00
Requirement already satisfied: idna in /usr/lib/python3/dist-packages (from kaggle) (3.3)
Collecting python-slugify
  Downloading python_slugify-8.0.4-py2.py3-none-any.whl (10 kB)
Collecting protobuf
  Downloading protobuf-6.30.2-cp39-abi3-manylinux2014_x86_64.whl (316 kB)
    316.2/316.2 KB 65.9 MB/s eta 0:00:00
Requirement already satisfied: setuptools>=21.0.0 in /usr/lib/python3/dist-packages (from kaggle) (59.6.0)
Collecting python-dateutil>=2.5.3
  Downloading python_dateutil-2.9.0.post0-py2.py3-none-any.whl (229 kB)
    229.9/229.9 KB 40.2 MB/s eta 0:00:00
Collecting charset-normalizer
  Downloading charset_normalizer-3.4.2-cp310-cp310-manylinux_2_17_x86_64_manylinux2014_x86_64.whl (149 kB)
    149.5/149.5 KB 39.2 MB/s eta 0:00:00
Requirement already satisfied: requests in /usr/lib/python3/dist-packages (from kaggle) (2.25.1)
Requirement already satisfied: certifi>=14.05.14 in /usr/lib/python3/dist-packages (from kaggle) (2020.6.20)
Requirement already satisfied: six>=1.10 in /usr/lib/python3/dist-packages (from kaggle) (1.16.0)
Collecting text-unidecode
  Downloading text_unidecode-1.3-py2.py3-none-any.whl (78 kB)
```

Kaggle CLI needs token authentication and correct permissions to work.

Create Kaggle Account and API Token

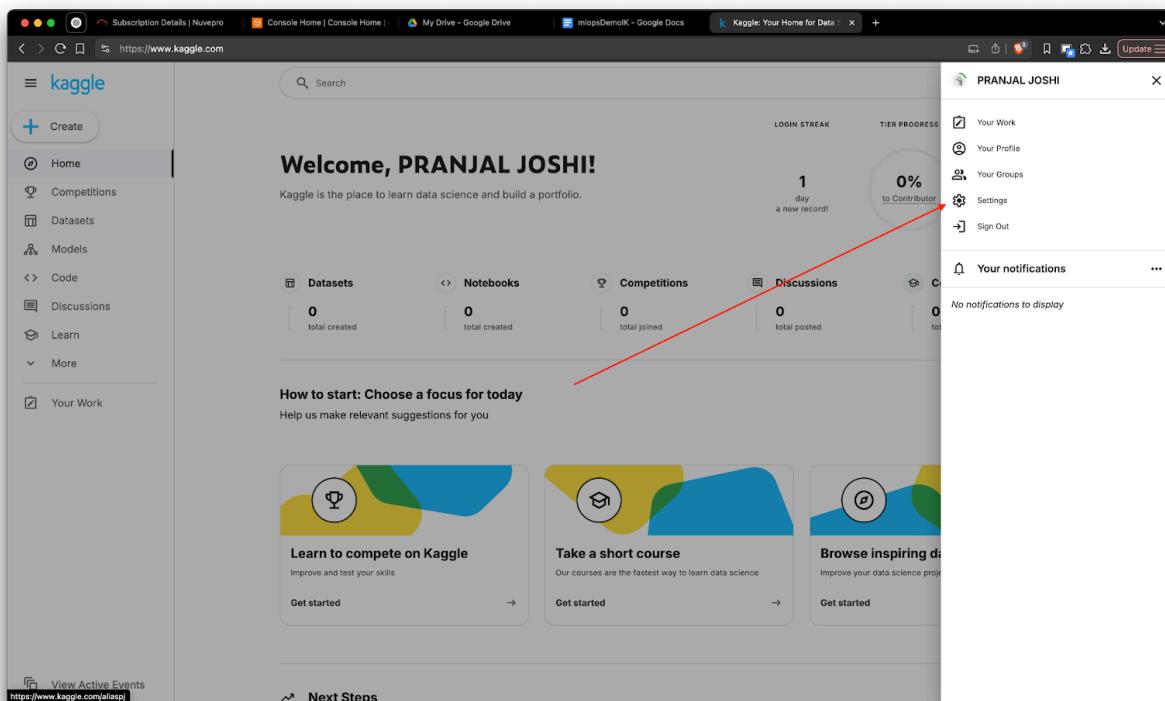
1. Go to <https://www.kaggle.com> and sign up or log in.
2. Click on your profile picture (top-right) > "Account".
3. Scroll down to the **API** section.
4. Click **"Create New API Token"**. This downloads a `kaggle.json` file.

5. Upload this to your EC2 instance under `~/kaggle/` as described above.

1. Join competition: [IEEE Fraud Detection](#)
2. On EC2:

linux commands:

```
→ kaggle competitions download -c ieee-fraud-detection  
→ unzip ieee-fraud-detection.zip -d ieee_fraud  
→ aws s3 cp ieee_fraud/ s3://fraud-demo-data/raw/ --recursive  
→ aws s3 ls s3://fraud-demo-data/raw/
```



The screenshot shows the 'Settings' page on the Kaggle website. On the left, there's a sidebar with links like 'Create', 'Home', 'Competitions', 'Datasets', 'Models', 'Code', 'Discussions', 'Learn', and 'More'. Below that is a 'Your Work' section with a 'View Active Events' button. The main content area has a search bar at the top. Under 'Settings', there are tabs for 'Account' and 'Notifications', with 'Account' being the active tab. It includes a 'Verify my account' button and a 'Theme' section with a dropdown set to 'Light theme'. The 'API' section contains text about Kaggle's beta API and two buttons: 'Create New Token' (highlighted with a red arrow) and 'Expire Token'. Below this is a 'Quotas' section with tables for 'Private Datasets', 'Private Models', 'Kaggle GPU', and 'Kaggle TPU'. At the bottom is an 'Active logins' section.

This screenshot is identical to the one above, but it includes a green checkmark icon and the text 'Ensure kaggle.json is in the location ~/kaggle/kaggle.json to use the API.' next to the 'Create New Token' button. A small 'Dismiss' link is also visible next to this message.

Kaggle

IEEE-CIS Fraud Detection

sample_submission.csv (6.08 MB)

Competition Rules

To see this data you need to agree to the [competition rules](#).
Join the competition to view the data.

Join the competition

Data Explorer

1.35 GB

- sample_submission.csv
- test_identity.csv
- test_transaction.csv
- train_identity.csv
- train_transaction.csv

Summary

- 5 files
- 871 columns

Download All

Metadata

License

View Active Events

```

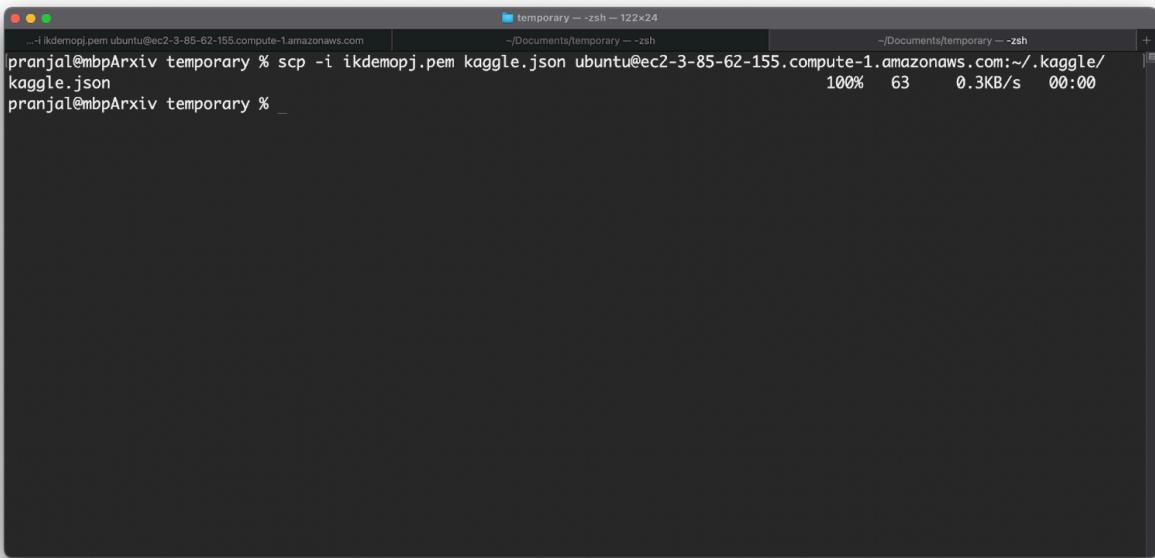
temporary -- ubuntu@ip-172-31-82-254: ~ ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com -- 183x38
...ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com ~Documents/temporary -- zsh
Requirement already satisfied: tqdm in ./local/lib/python3.10/site-packages (from kaggle) (4.67.1)
Requirement already satisfied: charset-normalizer in ./local/lib/python3.10/site-packages (from kaggle) (3.4.2)
Requirement already satisfied: urllib3<1.25.1 in /usr/lib/python3/dist-packages (from kaggle) (1.26.5)
Requirement already satisfied: requests in /usr/lib/python3/dist-packages (from kaggle) (2.25.1)
Requirement already satisfied: python-dateutil<2.5.3 in ./local/lib/python3.10/site-packages (from kaggle) (2.9.0.post0)
Requirement already satisfied: idna in /usr/lib/python3/dist-packages (from kaggle) (3.3)
Requirement already satisfied: python-slugify in ./local/lib/python3.10/site-packages (from kaggle) (8.0.4)
Requirement already satisfied: webencodings in ./local/lib/python3.10/site-packages (from kaggle) (0.5.1)
Requirement already satisfied: certifi>=2022.0.1 in ./local/lib/python3.10/site-packages (from kaggle) (30.2)
Requirement already satisfied: bleach in ./local/lib/python3.10/site-packages (from kaggle) (6.2.0)
Requirement already satisfied: s3transfer<0.13.0,>=0.12.0 in ./local/lib/python3.10/site-packages (from boto3) (0.12.0)
Requirement already satisfied: botocore<1.39.0,>=1.38.14 in ./local/lib/python3.10/site-packages (from boto3) (1.38.14)
Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in ./local/lib/python3.10/site-packages (from boto3) (1.0.1)
Requirement already satisfied: numpy>=1.22.4 in ./local/lib/python3.10/site-packages (from pandas) (2.2.5)
Requirement already satisfied: tzdata<2022.7 in ./local/lib/python3.10/site-packages (from pandas) (2025.2)
Requirement already satisfied: pytz=>2021.1 in /usr/lib/python3/dist-packages (from pandas) (2022.1)
ubuntu@ip-172-31-82-254: $ kaggle competitions download -c ieee-fraud-detection
kaggle: command not found
ubuntu@ip-172-31-82-254: $ pip3 install kaggle
echo 'export PATH="$PATH:./local/bin"' >> .bashrc
source ~/.bashrc
Defaulting to user installation because normal site-packages is not writable
Requirement already satisfied: kaggle in ./local/lib/python3.10/site-packages (1.7.4.5)
Requirement already satisfied: tabulate in ./local/lib/python3.10/site-packages (from kaggle) (4.07.1)
Requirement already satisfied: idna<3.1,!=3.0.0 in ./local/lib/python3.10/site-packages (from kaggle) (2.6.5)
Requirement already satisfied: certifi>=2022.0.1 in /usr/lib/python3/dist-packages (from kaggle) (2020.6.20)
Requirement already satisfied: six<1.10 in /usr/lib/python3/dist-packages (from kaggle) (1.16.0)
Requirement already satisfied: setuptools>=21.0.0 in /usr/lib/python3/dist-packages (from kaggle) (59.6.0)
Requirement already satisfied: requests in /usr/lib/python3/dist-packages (from kaggle) (2.25.1)
Requirement already satisfied: protobuf in ./local/lib/python3.10/site-packages (from kaggle) (6.30.2)
Requirement already satisfied: charset-normalizer in ./local/lib/python3.10/site-packages (from kaggle) (3.4.2)
Requirement already satisfied: python-slugify in ./local/lib/python3.10/site-packages (from kaggle) (8.0.4)
Requirement already satisfied: webencodings in ./local/lib/python3.10/site-packages (from kaggle) (0.5.1)
Requirement already satisfied: python-dateutil<2.5.3 in ./local/lib/python3.10/site-packages (from kaggle) (2.9.0.post0)
Requirement already satisfied: idna in /usr/lib/python3/dist-packages (from kaggle) (3.3)
Requirement already satisfied: text-unidecode in ./local/lib/python3.10/site-packages (from kaggle) (1.3)
ubuntu@ip-172-31-82-254: $
```

```
temporary — ubuntu@ip-172-31-82-254: ~ — ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com — 183x38
...-i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com ~Documents/temporary — zsh
[ubuntu@ip-172-31-82-254: ~]$ kaggle competitions download -c ieee-fraud-detection
403 Client Error: Forbidden for url: https://www.kaggle.com/api/v1/competitions/data/download-all/ieee-fraud-detection
[ubuntu@ip-172-31-82-254: ~]$ kaggle competitions download -c ieee-fraud-detection
Downloading ieee-fraud-detection.zip to /home/ubuntu
93%|██████████| 110M/118M [00:00<00:00, 1.15GB/s]
100%|██████████| 118M/118M [00:00<00:00, 1.09GB/s]
ubuntu@ip-172-31-82-254: ~$
```

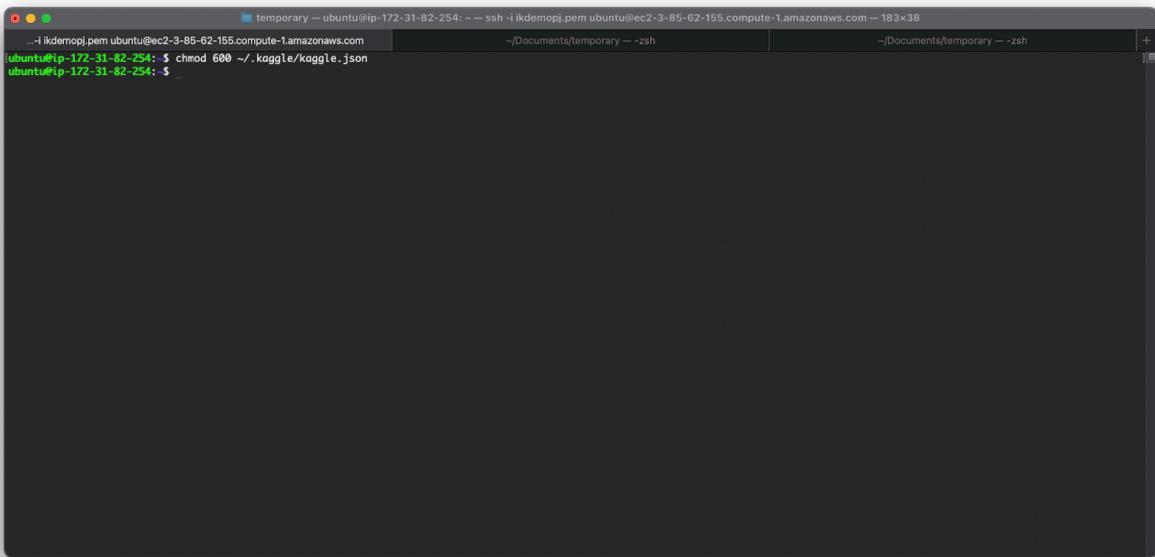
```
temporary — ubuntu@ip-172-31-82-254: ~ — ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com — 123x44
~/Documents/temporary — ubuntu@ip-172-31-82-254: ~ — ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com
[ubuntu@ip-172-31-82-254: ~]$ mkdir ~/.kaggle
[ubuntu@ip-172-31-82-254: ~]$
```

(

(in local windows/linux/mac terminal)



```
...-i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com temporary -zsh 122x24
pranjal@mbpArxiv temporary % scp -i ikdemopj.pem kaggle.json ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com:~/kaggle/
kaggle.json                                         100%   63    0.3KB/s  00:00
pranjal@mbpArxiv temporary %
```



```
...-i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com temporary -zsh 183x38
ubuntu@ip-172-31-82-254: ~ -- ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com -- 183x38
[ubuntu@ip-172-31-82-254: ~] $ chmod 600 ~/kaggle/kaggle.json
[ubuntu@ip-172-31-82-254: ~]
```

```
temporary -- ubuntu@ip-172-31-82-254: ~ ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com -- 183x38
...-i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com ~Documents/temporary --zsh
[ubuntu@ip-172-31-82-254: ~ kaggle competitions download -c ieee-fraud-detection
403 Client Error: Forbidden for url: https://www.kaggle.com/api/v1/competitions/data/download-all/ieee-fraud-detection
[ubuntu@ip-172-31-82-254: ~ kaggle competitions download -c ieee-fraud-detection
Downloading ieee-fraud-detection.zip to /home/ubuntu
93%|██████████| 118M/118M [00:00<00:00, 1.15GB/s]
100%|██████████| 118M/118M [00:00<00:00, 1.09GB/s]
[ubuntu@ip-172-31-82-254: ~ unzip ieee-fraud-detection.zip -d ieee_fraud
Archive: ieee_fraud-detection.zip
inflecting: ieee_Fraud/sample_submission.csv
inflecting: ieee_Fraud/test_identity.csv
inflecting: ieee_Fraud/test_transaction.csv
inflecting: ieee_Fraud/train_identity.csv
inflecting: ieee_Fraud/train_transaction.csv
[ubuntu@ip-172-31-82-254: ~
```

```
temporary -- ubuntu@ip-172-31-82-254: ~ ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com -- 183x38
...-i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com ~Documents/temporary --zsh
[ubuntu@ip-172-31-82-254: ~ python3 pycscripts3.py
Traceback (most recent call last):
  File "/home/ubuntu/pyscripts3.py", line 3, in <module>
    obj = s3.get_object(Bucket='fraud-demo-data-pj', Key='raw/train_transaction.csv')
  File "/home/ubuntu/.local/lib/python3.10/site-packages/botocore/client.py", line 570, in _api_call
    return self._make_api_call(operation_name, kwargs)
  File "/home/ubuntu/.local/lib/python3.10/site-packages/botocore/context.py", line 123, in wrapper
    return func(*args, **kwargs)
  File "/home/ubuntu/.local/lib/python3.10/site-packages/botocore/client.py", line 1031, in _make_api_call
    raise error_class(parsed_response, operation_name)
botocore.errorfactory.NoSuchKey: An error occurred (NoSuchKey) when calling the GetObject operation: The specified key does not exist.
[ubuntu@ip-172-31-82-254: ~ ls
ieee_fraud-detection.zip  ieee_fraud  pycscripts3.py
[ubuntu@ip-172-31-82-254: ~ cd ieee_Fraud/
[ubuntu@ip-172-31-82-254: ~/ieee_Fraud$ :
-bash: syntax error near unexpected token `;'
[ubuntu@ip-172-31-82-254: ~/ieee_Fraud$ ls
sample_submission.csv  test_identity.csv  test_transaction.csv  train_identity.csv  train_transaction.csv
[ubuntu@ip-172-31-82-254: ~/ieee_Fraud$ aws s3 ls s3://fraud-demo-data-pj/raw/
[ubuntu@ip-172-31-82-254: ~/ieee_Fraud$ cd ..
[ubuntu@ip-172-31-82-254: ~ ls
ieee_fraud-detection.zip  ieee_fraud  pycscripts3.py
[ubuntu@ip-172-31-82-254: ~ aws s3 ls s3://fraud-demo-data-pj/raw/
[ubuntu@ip-172-31-82-254: ~ aws s3 cp ieee_fraud/train_transaction.csv s3://fraud-demo-data-pj/raw/
upload: ieee_fraud/train_transaction.csv to s3://fraud-demo-data-pj/raw/train_transaction.csv
[ubuntu@ip-172-31-82-254: ~
```

Simulates a typical pipeline where ingestion uploads raw data to the cloud.

17 Step 5: Basic Data Ingestion + Logging Demo

Run this command to specify region for cloudwatch:

→**export AWS_DEFAULT_REGION=us-east-1**

Run this script using editor by running command:

→vi script.py

[Guide to vi editor](#)

```
#script.py

import boto3, pandas as pd, logging, watchtower

# Fetch from S3

s3 = boto3.client('s3')

obj = s3.get_object(Bucket='fraud-demo-data-<yourname>',
Key='raw/train_transaction.csv')

df = pd.read_csv(obj['Body'], nrows=5)

print(df)

# CloudWatch logging with region_name

logger = logging.getLogger(__name__)

logger.setLevel(logging.INFO)

logger.addHandler(watchtower.CloudWatchLogHandler(log_group='fraud-de
mo-logs'))

logger.info("Loaded 5 rows from S3 successfully")
```

```
temporary -- ubuntu@ip-172-31-82-254: ~ ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com -- 183x38
...-i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com
import boto3, pandas as pd
s3 = boto3.client('s3')
obj = s3.get_object(Bucket='fraud-demo-data', Key='raw/train_transaction.csv')
df = pd.read_csv(obj['Body'])
print(df.head())
~
~
~
```

```
temporary -- ubuntu@ip-172-31-82-254: ~ ssh -i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com -- 183x38
...-i ikdemopj.pem ubuntu@ec2-3-85-62-155.compute-1.amazonaws.com
~/Documents/temporary -- zsh
~/Documents/temporary -- zsh
Scanning candidates...
Scanning linux images...
Running kernel seems to be up-to-date.

Restarting services...
Service restarts being deferred:
/etc/needrestart/restart.d/dbus.service
systemctl restart networkd-dispatcher.service
systemctl restart polkit.service
systemctl restart unattended-upgrades.service

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
[ubuntu@ip-172-31-82-254: ~]
[ubuntu@ip-172-31-82-254: ~]
[ubuntu@ip-172-31-82-254: ~] ls
ieee-Fraud-detection.zip  ieee_Fraud
[ubuntu@ip-172-31-82-254: ~] pip3 install watchtower
Defaulting to user installation because normal site-packages is not writable
Collecting watchtower
  Downloading watchtower-3.4.0-py3-none-any.whl (18 kB)
Requirement already satisfied: boto3<2,>=1.9.253 in ./local/lib/python3.10/site-packages (from watchtower) (1.38.14)
Requirement already satisfied: s3transfer<0.13.0,>=0.12.0 in ./local/lib/python3.10/site-packages (from boto3<2,>=1.9.253->watchtower) (0.12.0)
Requirement already satisfied: botocore<1.39.0,>=1.38.14 in ./local/lib/python3.10/site-packages (from boto3<2,>=1.9.253->watchtower) (1.38.14)
Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in ./local/lib/python3.10/site-packages (from boto3<2,>=1.9.253->watchtower) (1.0.1)
Requirement already satisfied: urllib3!=2.2.0,>=1.25.4 in /usr/lib/python3/dist-packages (from botocore<1.39.0,>=1.38.14->boto3<2,>=1.9.253->watchtower) (1.26.5)
Requirement already satisfied: python-dateutil<3.0.0,>=2.1 in ./local/lib/python3.10/site-packages (from botocore<1.39.0,>=1.38.14->boto3<2,>=1.9.253->watchtower) (2.9.0.post0)
Requirement already satisfied: six<1.5 in /usr/lib/python3/dist-packages (from python-dateutil<3.0.0,>=2.1->botocore<1.39.0,>=1.38.14->boto3<2,>=1.9.253->watchtower) (1.16.0)
Installing collected packages: watchtower
Successfully installed watchtower-3.4.0
[ubuntu@ip-172-31-82-254: ~] vi pyscripts3.py
[ubuntu@ip-172-31-82-254: ~] ls
ieee-Fraud-detection.zip  ieee_Fraud  pyscripts3.py
[ubuntu@ip-172-31-82-254: ~] python3 pyscripts3.py
```

```
temporary — ubuntu@ip-172-31-92-42: ~ — ssh -i ikdemopj.pem ubuntu@ec2-44-210-117-176.compute-1.amazonaws.com — 111x27
...-i ikdemopj.pem ubuntu@ec2-44-210-117-176.compute-1.amazonaws.com          ~/Documents/temporary — zsh
ubuntu@ip-172-31-92-42:~$ export AWS_DEFAULT_REGION=us-east-1
ubuntu@ip-172-31-92-42:~$ vi scriptChunks.py
ubuntu@ip-172-31-92-42:~$ python3 scriptChunks.py
   TransactionID  isFraud  TransactionDT  TransactionAmt  ...  V336  V337  V338  V339
0      2987000       0        86400       68.5  ...  NaN  NaN  NaN  NaN
1      2987001       0        86401       29.0  ...  NaN  NaN  NaN  NaN
2      2987002       0        86469       59.0  ...  NaN  NaN  NaN  NaN
3      2987003       0        86499       50.0  ...  NaN  NaN  NaN  NaN
4      2987004       0        86506       50.0  ...  0.0  0.0  0.0  0.0
[5 rows x 394 columns]
ubuntu@ip-172-31-92-42:~$ cat scriptChunks.py
import boto3, pandas as pd, logging, watchtower

# Fetch from S3
s3 = boto3.client('s3')
obj = s3.get_object(Bucket='fraud-demo-data-pj', Key='raw/train_transaction.csv')
df = pd.read_csv(obj['Body'], nrows=5)
print(df)

# CloudWatch logging with region_name
logger = logging.getLogger(__name__)
logger.setLevel(logging.INFO)
logger.addHandler(watchtower.CloudWatchLogHandler(log_group='fraud-demo-logs'))
logger.info("Loaded 5 rows from S3 successfully")

ubuntu@ip-172-31-92-42:~$
```

Output on the aws cloudwatch dashboard:

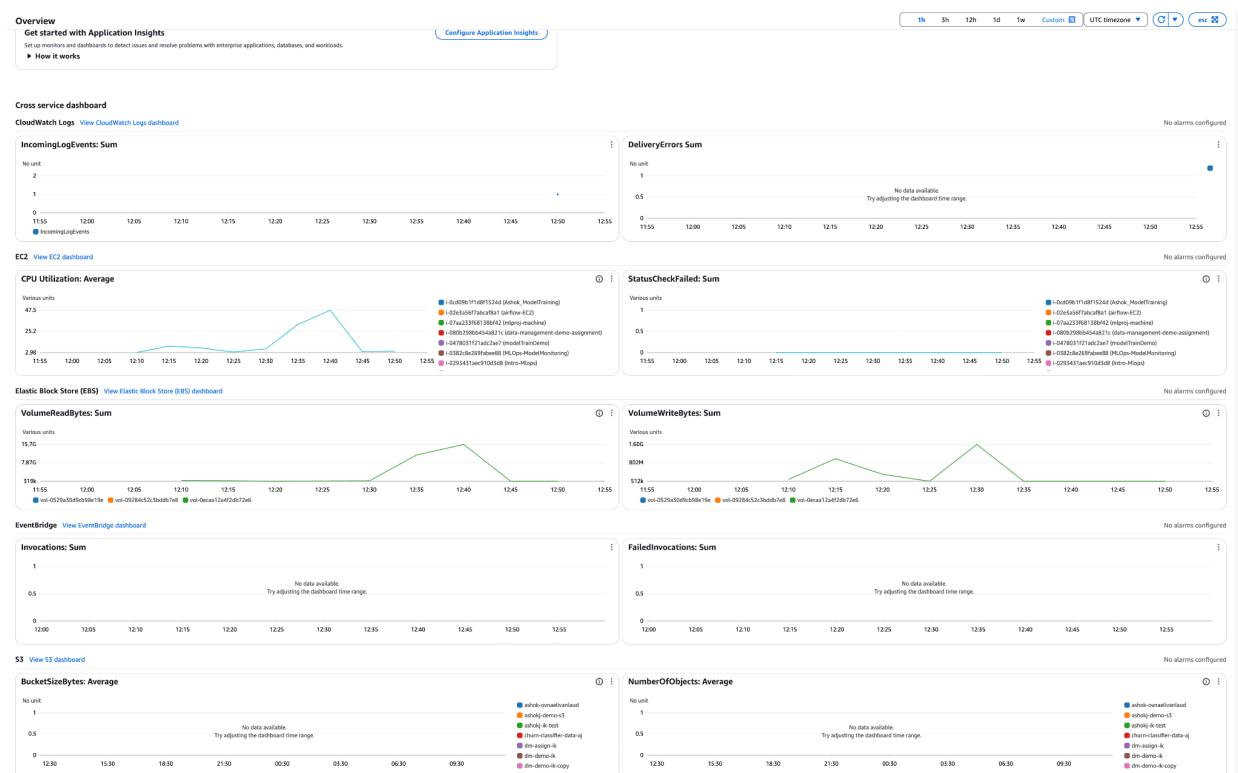
Screenshot of the AWS CloudWatch Home page (https://us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#home?:~:(timeRange=3600000))

The left sidebar contains links to various CloudWatch services: CloudWatch Metrics, CloudWatch Logs, CloudWatch Events, CloudWatch Application Insights, CloudWatch Network Metrics, CloudWatch Insights, CloudWatch Telemetry, CloudWatch Getting Started, and CloudWatch What's New.

The main content area features several sections:

- CloudWatch Services:** Includes links to CloudWatch Metrics, CloudWatch Logs, CloudWatch Events, CloudWatch Application Insights, CloudWatch Network Metrics, and CloudWatch Insights.
- CloudWatch Features:** Includes links to CloudWatch Metrics, CloudWatch Logs, CloudWatch Events, CloudWatch Application Insights, CloudWatch Network Metrics, and CloudWatch Insights.
- CloudWatch Documentation:** Includes links to Amazon Virtual Private Cloud (Amazon VPC) User Guide, Amazon Elastic Compute Cloud (EC2) User Guide, and Amazon Elastic Compute Cloud (EC2) Instance User Guide.
- CloudWatch Resources:** A search bar for introducing resource search.
- CloudWatch Knowledge Articles:** Includes links to CloudWatch Metrics, CloudWatch Logs, CloudWatch Events, CloudWatch Application Insights, CloudWatch Network Metrics, and CloudWatch Insights.

At the bottom, there are "Were these results helpful?" buttons (Yes/No), a CPU Utilization: Average chart, and a footer with copyright information.



Logging enables observability. Watchtower helps push logs directly to CloudWatch.

Common Errors

- AccessDenied: IAM role lacks permissions or typo in bucket/key
 - NoSuchKey: Wrong file name or missing upload
-

Bottlenecks and Engineering Solutions

- No dataset versioning → DVC, LakeFS
- No schema validation → Great Expectations, Pandera
- No retraining workflow → GitHub Actions, Jenkins
- Model not tracked → MLflow, SageMaker Model Registry
- No deployment pipeline → EKS, Lambda, SageMaker Endpoint
- No drift detection → EvidentlyAI, WhyLabs

These bottlenecks motivate the need for MLOps practices and tooling.

System Design Sketch

Raw Data (S3)



Validation + Feature Engineering (EC2/SageMaker)



Training (SageMaker)



Model Registry + CI/CD (MLflow/Jenkins)



Deployment (EKS / Lambda / SageMaker Endpoint)



Monitoring (CloudWatch + Drift Detector)

This architecture illustrates how production ML systems are modular, observable, and automated.
